

## Week 1

### Binary classifier

```
$ googleimagesdownload -kf sedan_versus_pickup.txt -o sedan_versus_pickup  
-l 100
```

manual cleanup: delete some images.

peach > things with the color peach, like hammocks or iphone cases

apple > things like apple pie, cans of apple juice etc

bmw> just the logo

also generally, some 0byte files / corrupted files

### Expectations

- giraffe\_versus\_elephant/
  - easy, very different shapes and also colors. 98%
- appel\_versus\_perzik/
  - tricky, very similar. perhaps 70%
- bmw\_versus\_mercedes/
  - may pick up on logos. 80%
- bmw\_z4\_versus\_bmw\_3/
  - easy, very different shapes. perhaps 90%
- chair\_versus\_table/
  - different shapes, but some odd photos. 80%
- meeuw\_versus\_spreeuw/
  - easy, very different shapes and also colors. 98%
- rutte\_versus\_harry\_potter/
  - tricky, dataset is not great, for harry potter there are some group pictures. but the colors from harry are more blueisch. Also rutte is always wearing a suit. I think 80% should be feasible
- sedan\_versus\_pickup/
  - very different shape, should reach 85%
- south\_african\_giraffe\_versus\_masai\_giraffe/
  - tricky one. I hope the NN is able to pick up on the different patterns.  
Maybe 60%

### # Uploading files

```
roland-macbook:fastai roland$ gcloud compute copy-files data gpu-docker-  
host-fromimage:~/data
```

WARNING: `gcloud compute copy-files` is deprecated. Please use `gcloud compute scp` instead. Note that `gcloud compute scp` does not have recursive copy on by default. To turn on recursion, use the `--recurse` flag.

```
# Split in train/valid
```

```
# Training model
```

```
OSErrror: Error handling image at: /fastai/data/data/giraffe_versus_elephant/
valid/elephant/56. elephant-loader.gif
```

I just removed this file. Perhaps gif is not a valid format.

Giraffe versus elephant

```
100% |██████████| 1/1 [00:01<00:00, 1.17s/it]
```

```
Epoch 100% 2/2 [00:02<00:00, 1.36s/it]
```

| epoch | trn_loss | val_loss | accuracy |
|-------|----------|----------|----------|
| 0     | 1.210398 | 0.370057 | 0.897436 |
| 1     | 0.731472 | 0.070873 | 1.0      |

```
# 1. A few correct labels at random
plot_val_with_title(rand_by_correct(True), "Correctly classified")
```

Correctly classified



```
In [69]: most_uncertain = np.argsort(np.abs(probs - 0.5))[:4]
plot_val_with_title(most_uncertain, "Most uncertain predictions")
```

Most uncertain predictions



Conclusion: my prediction of 98% accuracy was even too low! We achieved 100%.

BMW versus Mercedes

```
| In [121]: arch=resnet34  
data = ImageClassifierData.from_paths(PATH, tfms=tfms_from_model(arch, sz))  
learn = ConvLearner.pretrained(arch, data, precompute=True)  
learn.fit(0.01, 10)
```

Epoch  100% 10/10 [00:00<00:00, 18.84it/s]

| epoch | trn_loss | val_loss | accuracy |
|-------|----------|----------|----------|
| 0     | 0.877495 | 1.022667 | 0.513514 |
| 1     | 0.889549 | 0.852195 | 0.540541 |
| 2     | 0.771642 | 0.741717 | 0.513514 |
| 3     | 0.67758  | 0.674486 | 0.594595 |
| 4     | 0.611079 | 0.650615 | 0.675676 |
| 5     | 0.570702 | 0.647646 | 0.72973  |
| 6     | 0.529626 | 0.669904 | 0.756757 |
| 7     | 0.476845 | 0.701    | 0.756757 |
| 8     | 0.440523 | 0.72828  | 0.72973  |
| 9     | 0.421131 | 0.749718 | 0.756757 |

```
| Out[121]: [array([0.74972]), 0.7567567825317383]
```

Resultaat niet super. Gaat snel overfitten.

```
| In [119]: # 1. A few correct labels at random  
plot_val_with_title(rand_by_correct(True), "Correctly classified")
```

Correctly classified



```
| In [120]: # 2. A few incorrect labels at random  
plot_val_with_title(rand_by_correct(False), "Incorrectly classified")
```

Incorrectly classified



```
In [124]: plot_val_with_title(most_by_correct(0, True), "Most correct class 1")
```

Most correct class 1



```
In [125]: plot_val_with_title(most_by_correct(1, True), "Most correct class 2")
```

Most correct class 2



```
In [124]: plot_val_with_title(most_by_correct(0, True), "Most correct class 1")
```

Most correct class 1



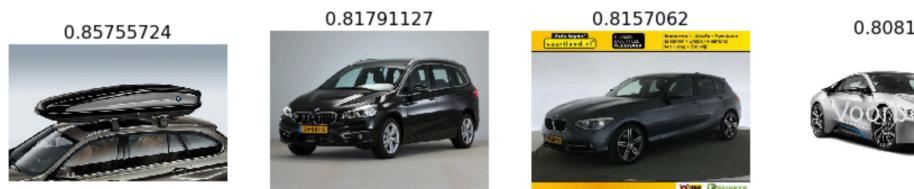
```
In [125]: plot_val_with_title(most_by_correct(1, True), "Most correct class 2")
```

Most correct class 2



```
In [126]: plot_val_with_title(most_by_correct(0, False), "Most incorrect class 1")
```

Most incorrect class 1



```
In [128]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")
```

Most incorrect class 2



```
In [126]: plot_val_with_title(most_by_correct(0, False), "Most incorrect class 1")
```

Most incorrect cats



```
In [128]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")
```

Most incorrect class 2



```
In [129]: most_uncertain = np.argsort(np.abs(probs - 0.5))[:4]
plot_val_with_title(most_uncertain, "Most uncertain predictions")
```

Most uncertain predictions



Conclusion: not great results. My intuition is that we need more data.

BMW 3 series versus BMW Z4

```
In [160]: arch=resnet34
data = ImageClassifierData.from_paths(PATH, tfms=tfms_from_model(arch, sz))
learn = ConvLearner.pretrained(arch, data, precompute=True)
learn.fit(0.01, 10)
```

```
100%|██████████| 3/3 [00:03<00:00,  1.02s/it]
100%|██████████| 1/1 [00:01<00:00,  1.02s/it]
```

Epoch   100% 10/10 [00:00<00:00, 18.56it/s]

| epoch | trn_loss | val_loss | accuracy |
|-------|----------|----------|----------|
| 0     | 0.872835 | 0.878617 | 0.512821 |
| 1     | 0.668088 | 0.475407 | 0.74359  |
| 2     | 0.53179  | 0.256914 | 0.923077 |
| 3     | 0.437537 | 0.15862  | 0.948718 |
| 4     | 0.365405 | 0.128076 | 0.948718 |
| 5     | 0.313121 | 0.119094 | 0.948718 |
| 6     | 0.268797 | 0.126063 | 0.923077 |
| 7     | 0.243382 | 0.142961 | 0.923077 |
| 8     | 0.216353 | 0.161994 | 0.948718 |
| 9     | 0.197005 | 0.176308 | 0.948718 |

```
Out[160]: [array([0.17631]), 0.9487179517745972]
```

```
In [172]: plot_val_with_title(most_by_correct(0, True), "Most correct class 1")
```

Most correct class 1



```
In [173]: plot_val_with_title(most_by_correct(1, True), "Most correct class 2")
```

Most correct class 2



```
In [172]: plot_val_with_title(most_by_correct(0, True), "Most correct class 1")
```

Most correct class 1



```
In [173]: plot_val_with_title(most_by_correct(1, True), "Most correct class 2")
```

Most correct class 2



```
In [174]: plot_val_with_title(most_by_correct(0, False), "Most incorrect class 1")
Most incorrect class 1
<Figure size 1152x576 with 0 Axes>
```

```
In [175]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")
Most incorrect class 2
```



```
In [174]: plot_val_with_title(most_by_correct(0, False), "Most incorrect class 1")
Most incorrect class 1
<Figure size 1152x576 with 0 Axes>
```

```
In [175]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")
Most incorrect class 2
```



```
In [176]: most_uncertain = np.argsort(np.abs(probs - 0.5))[:4]
plot_val_with_title(most_uncertain, "Most uncertain predictions")
```

Most uncertain predictions



Conclusion: as expected, these classes are relatively easy to distinguish.

CHAIR versus TABLE

```
In [198]: arch=resnet34
data = ImageClassifierData.from_paths(PATH, tfms=tfms_from_model(arch, sz))
learn = ConvLearner.pretrained(arch, data, precompute=True)
learn.fit(0.01, 10)
```

100% |████████| 1/1 [00:00<00:00, 1.06it/s]

Epoch  100% 10/10 [00:00<00:00, 20.44it/s]

| epoch | trn_loss | val_loss | accuracy |
|-------|----------|----------|----------|
| 0     | 0.563126 | 0.944147 | 0.529412 |
| 1     | 0.524812 | 0.529544 | 0.676471 |
| 2     | 0.534502 | 0.329214 | 0.852941 |
| 3     | 0.414387 | 0.198126 | 0.941176 |
| 4     | 0.347168 | 0.120709 | 0.970588 |
| 5     | 0.292543 | 0.090168 | 1.0      |
| 6     | 0.251866 | 0.072425 | 1.0      |
| 7     | 0.220074 | 0.059451 | 1.0      |
| 8     | 0.195416 | 0.055396 | 1.0      |
| 9     | 0.199023 | 0.055025 | 0.970588 |

```
Out[198]: [array([0.05503]), 0.970588207244873]
```

```
In [207]: # 1. A few correct labels at random
plot_val_with_title(rand_by_correct(True), "Correctly classified")
```

Correctly classified

0.9999685



0.003709444

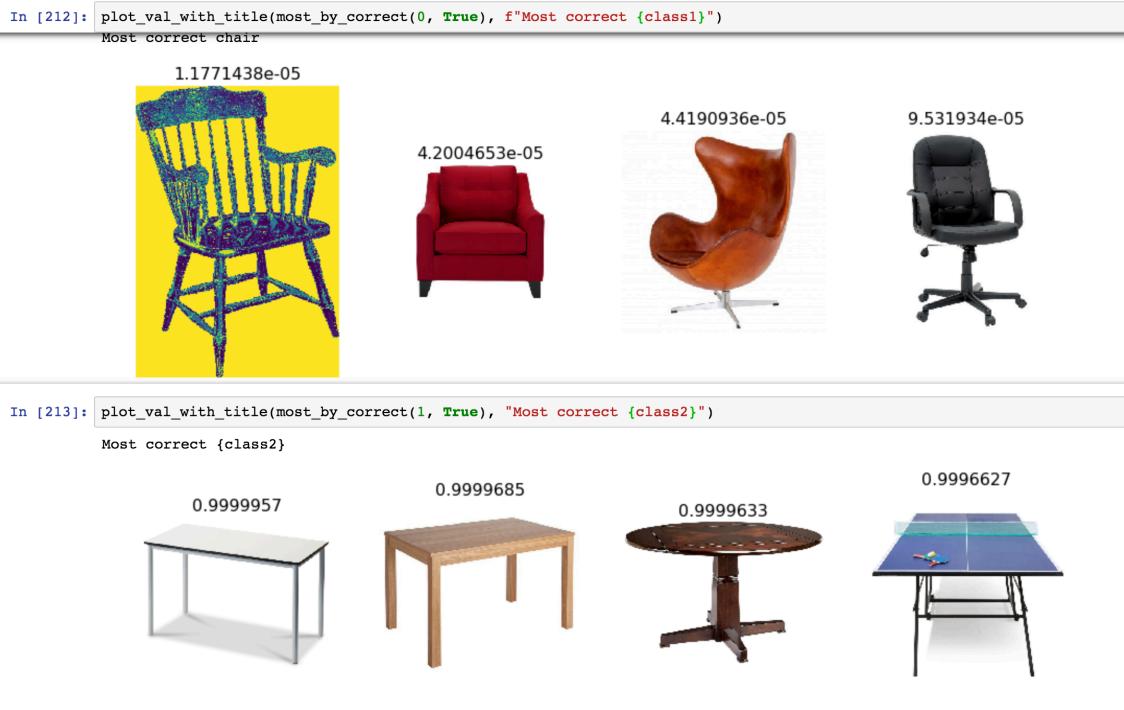


0.0024127823



0.95282674





► In [215]: `plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")`

Most incorrect class 2

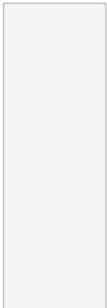
0.38439244



```
In [215]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")  
Most incorrect class 2  
0.38439244
```



```
In [216]: most_uncertain = np.argsort(np.abs(probs -0.5))[:4]  
plot_val_with_title(most_uncertain, "Most uncertain predictions")  
Most uncertain predictions  
0.38439244 0.7602744 0.12132331 0.11358159
```



Conclusion: does much better than my predicted 80%! Cool.

**meeuw\_versus\_spreeuw**

```
In [229]: arch=resnet34
data = ImageClassifierData.from_paths(PATH, tfms=tfms_from_model(arch, sz))
learn = ConvLearner.pretrained(arch, data, precompute=True)
learn.fit(0.01, 10)
```

100%|██████████| 3/3 [00:03<00:00, 1.23s/it]  
100%|██████████| 1/1 [00:02<00:00, 2.32s/it]

Epoch 100% 10/10 [00:00<00:00, 17.17it/s]

| epoch | trn_loss | val_loss | accuracy |
|-------|----------|----------|----------|
| 0     | 0.709314 | 0.528505 | 0.74359  |
| 1     | 0.548378 | 0.233248 | 1.0      |
| 2     | 0.410591 | 0.130629 | 0.974359 |
| 3     | 0.321861 | 0.094448 | 0.974359 |
| 4     | 0.268481 | 0.081766 | 0.974359 |
| 5     | 0.234602 | 0.085104 | 0.974359 |
| 6     | 0.200543 | 0.089793 | 0.974359 |
| 7     | 0.173074 | 0.095044 | 0.974359 |
| 8     | 0.150757 | 0.094129 | 0.974359 |
| 9     | 0.139575 | 0.093612 | 0.974359 |

Out[229]: [array([0.09361]), 0.9743589758872986]

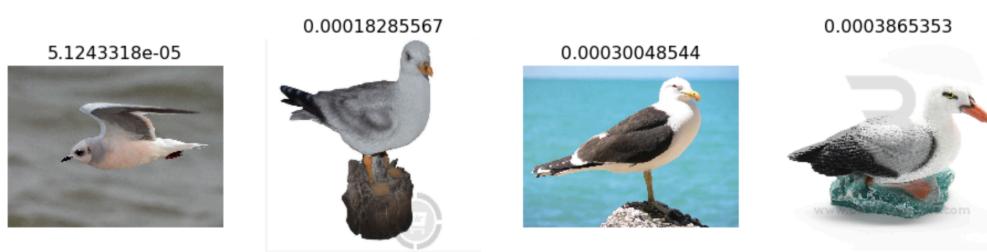
```
In [238]: # 1. A few correct labels at random
plot_val_with_title(rand_by_correct(True), "Correctly classified")
```

Correctly classified



```
In [241]: plot_val_with_title(most_by_correct(0, True), "Most correct {class1}")
```

Most correct meeuw



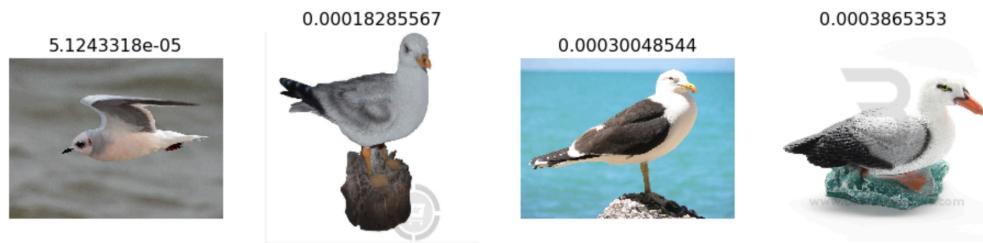
```
In [242]: plot_val_with_title(most_by_correct(1, True), "Most correct {class2}")
```

Most correct {class2}



```
In [241]: plot_val_with_title(most_by_correct(0, True), f"Most correct {class1}")
```

Most correct meeuw



```
In [242]: plot_val_with_title(most_by_correct(1, True), "Most correct {class2}")
```

Most correct {class2}



```
In [244]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")
```

Most incorrect class 2

0.07857431



```
In [244]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")
```

```
Most incorrect class 2
```

```
0.07857431
```



Conclusion: this dataset can be predicted like I expected. High accuracy given the very different classes.

## SEDAN versus Pickup truck

```
In [265]: arch=resnet34
data = ImageClassifierData.from_paths(PATH, tfms=tfms_from_model(arch, sz))
learn = ConvLearner.pretrained(arch, data, precompute=True)
learn.fit(0.01, 15)
```

```
100% |██████████| 3/3 [00:04<00:00, 1.39s/it]
100% |██████████| 1/1 [00:01<00:00, 1.37s/it]
```

```
Epoch [██████████] 100% 15/15 [00:00<00:00, 18.65it/s]
```

| epoch | trn_loss | val_loss | accuracy |
|-------|----------|----------|----------|
| 0     | 0.836447 | 0.466323 | 0.868421 |
| 1     | 0.643211 | 0.248251 | 0.947368 |
| 2     | 0.473987 | 0.175111 | 0.947368 |
| 3     | 0.410001 | 0.155383 | 0.947368 |
| 4     | 0.338402 | 0.14888  | 0.947368 |
| 5     | 0.289612 | 0.150993 | 0.947368 |
| 6     | 0.253184 | 0.151374 | 0.947368 |
| 7     | 0.221403 | 0.151549 | 0.947368 |
| 8     | 0.198052 | 0.152612 | 0.947368 |
| 9     | 0.177162 | 0.150671 | 0.947368 |
| 10    | 0.163218 | 0.148143 | 0.947368 |
| 11    | 0.154793 | 0.142629 | 0.947368 |
| 12    | 0.140006 | 0.140853 | 0.947368 |

```
In [278]: plot_val_with_title(most_by_correct(0, True), f"Most correct {class1}")
```

Most correct pickup\_truck



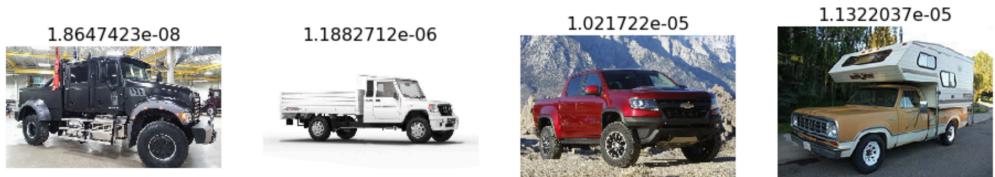
```
In [279]: plot_val_with_title(most_by_correct(1, True), "Most correct {class2}")
```

Most correct {class2}



```
In [278]: plot_val_with_title(most_by_correct(0, True), f"Most correct {class1}")
```

Most correct pickup\_truck



```
In [279]: plot_val_with_title(most_by_correct(1, True), "Most correct {class2}")
```

Most correct {class2}



```
In [280]: plot_val_with_title(most_by_correct(0, False), "Most incorrect class 1")
```



```
► In [281]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")
```

Most incorrect class 2

0.07814792



```
In [281]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")
```

Most incorrect class 2

0.07814792



I think the “most incorrect” are actually correct! its just that the labels are wrong!

```
In [282]: most_uncertain = np.argsort(np.abs(probs - 0.5))[:4]
plot_val_with_title(most_uncertain, "Most uncertain predictions")
```



These are fairly small, vague pictures.

Conclusion: 85% was easily achieved.

Rutte versus Harry Potter

```
In [295]: arch=resnet34
data = ImageClassifierData.from_paths(PATH, tfms=tfms_from_model(arch, sz))
learn = ConvLearner.pretrained(arch, data, precompute=True)
learn.fit(0.01, 10)
```

```
100%|██████████| 3/3 [00:02<00:00,  1.03it/s]
100%|██████████| 1/1 [00:00<00:00,  1.29it/s]
```

```
Epoch 100% 10/10 [00:00<00:00, 19.27it/s]
```

| epoch | trn_loss | val_loss | accuracy |
|-------|----------|----------|----------|
| 0     | 0.710575 | 0.479968 | 0.8      |
| 1     | 0.510306 | 0.270477 | 0.942857 |
| 2     | 0.392461 | 0.146947 | 1.0      |
| 3     | 0.310633 | 0.117591 | 0.942857 |
| 4     | 0.261007 | 0.098725 | 0.971429 |
| 5     | 0.363741 | 0.106628 | 0.971429 |
| 6     | 0.308863 | 0.082998 | 0.971429 |
| 7     | 0.307874 | 0.061038 | 0.971429 |
| 8     | 0.391125 | 0.076511 | 0.971429 |
| 9     | 0.346723 | 0.100858 | 0.942857 |

```
Out[295]: [array([0.10086]), 0.9428571462631226]
```

```
In [304]: # 1. A few correct labels at random
plot_val_with_title(rand_by_correct(True), "Correctly classified")
```

Correctly classified



```
In [304]: # 1. A few correct labels at random
plot_val_with_title(rand_by_correct(True), "Correctly classified")
```

Correctly classified



```
In [310]: # 2. A few incorrect labels at random  
plot_val_with_title(rand_by_correct(False), "Incorrectly classified")
```

Incorrectly classified

0.58405894



0.6912566

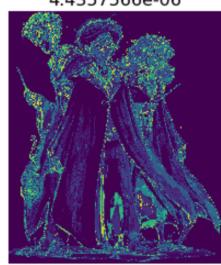


0.6



```
In [311]: plot_val_with_title(most_by_correct(0, True), f"Most correct {class1}")
```

4.4357566e-06



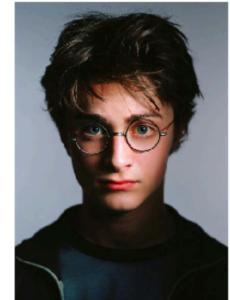
0.00037412



0.0016213414

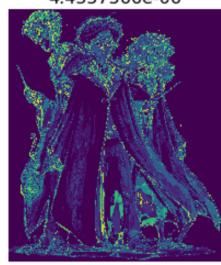


0.008630459



```
In [311]: plot_val_with_title(most_by_correct(0, True), f"Most correct {class1}")
```

4.4357566e-06



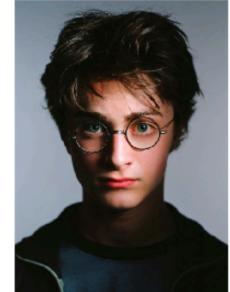
0.00037412



0.0016213414



0.008630459



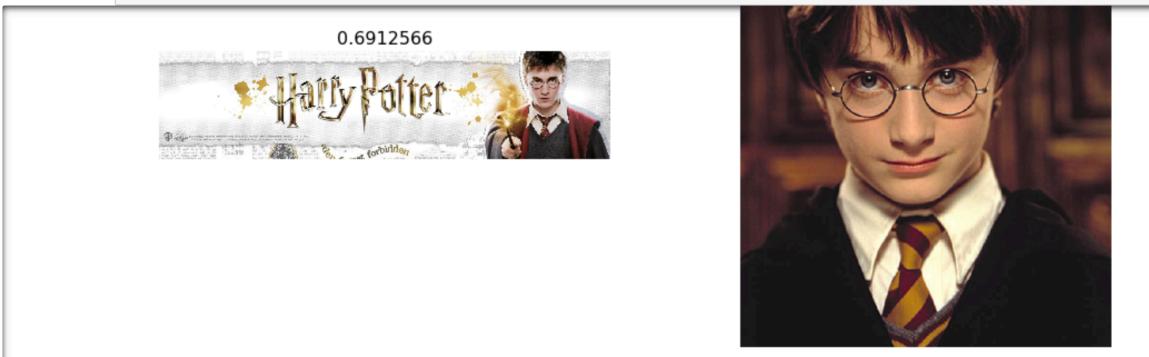
```
In [311]: plot_val_with_title(most_by_correct(0, True), f"Most correct {class1}")
```



```
In [312]: plot_val_with_title(most_by_correct(1, True), "Most correct {class2}")
```



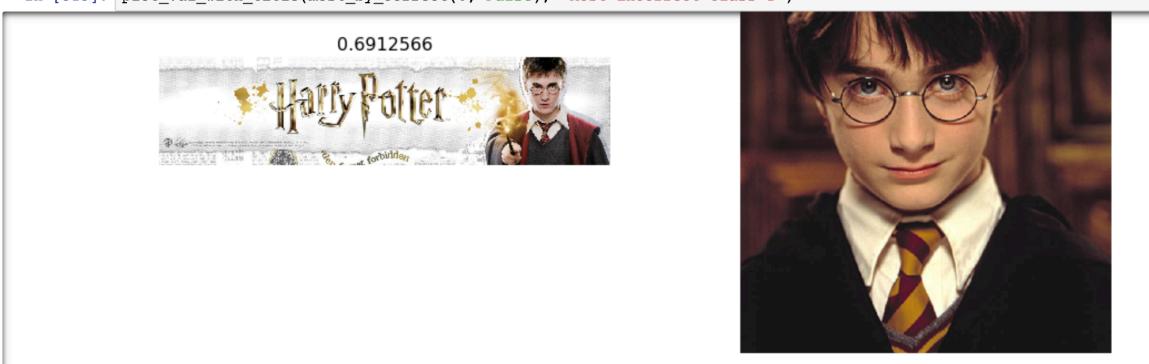
```
In [313]: plot_val_with_title(most_by_correct(0, False), "Most incorrect class 1")
```



```
In [314]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")
```

```
Most incorrect class 2  
<Figure size 1152x576 with 0 Axes>
```

```
In [313]: plot_val_with_title(most_by_correct(0, False), "Most incorrect class 1")
```



```
In [314]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")
```

```
Most incorrect class 2  
<Figure size 1152x576 with 0 Axes>
```

```
In [315]: most_uncertain = np.argsort(np.abs(probs - 0.5))[:4]
plot_val_with_title(most_uncertain, "Most uncertain predictions")
```



```
In [315]: most_uncertain = np.argsort(np.abs(probs - 0.5))[:4]
plot_val_with_title(most_uncertain, "Most uncertain predictions")
```



Conclusion: again, the NN performed much better than I expected. Though I am not sure whether it can really recognize these 2 persons. I think the style of the photos is just really different. Mental note: do a comparison of various politicians or actors.

Last: South African Giraffe versus Masai Giraffe

```
In [327]: arch=resnet34
data = ImageClassifierData.from_paths(PATH, tfms=tfms_from_model(arch, sz))
learn = ConvLearner.pretrained(arch, data, precompute=True)
learn.fit(0.01, 10)
```

100% |██████████| 3/3 [00:05<00:00, 1.70s/it]  
100% |██████████| 1/1 [00:01<00:00, 1.96s/it]

Epoch  100% 10/10 [00:00<00:00, 17.28it/s]

| epoch | trn_loss | val_loss | accuracy |
|-------|----------|----------|----------|
| 0     | 0.755909 | 0.806627 | 0.45     |
| 1     | 0.729105 | 0.759772 | 0.5      |
| 2     | 0.684765 | 0.71894  | 0.575    |
| 3     | 0.633249 | 0.695716 | 0.55     |
| 4     | 0.580498 | 0.718255 | 0.575    |
| 5     | 0.51721  | 0.747415 | 0.55     |
| 6     | 0.478577 | 0.779054 | 0.575    |
| 7     | 0.447871 | 0.822055 | 0.575    |
| 8     | 0.409474 | 0.847211 | 0.575    |
| 9     | 0.376428 | 0.862352 | 0.575    |

```
Out[327]: [array([0.86235]), 0.574999988079071]
```

```
In [327]: arch=resnet34  
data = ImageClassifierData.from_paths(PATH, tfms=tfms_from_model(arch, sz))  
learn = ConvLearner.pretrained(arch, data, precompute=True)  
learn.fit(0.01, 10)
```

```
100%|██████████| 3/3 [00:05<00:00, 1.70s/it]  
100%|██████████| 1/1 [00:01<00:00, 1.96s/it]
```

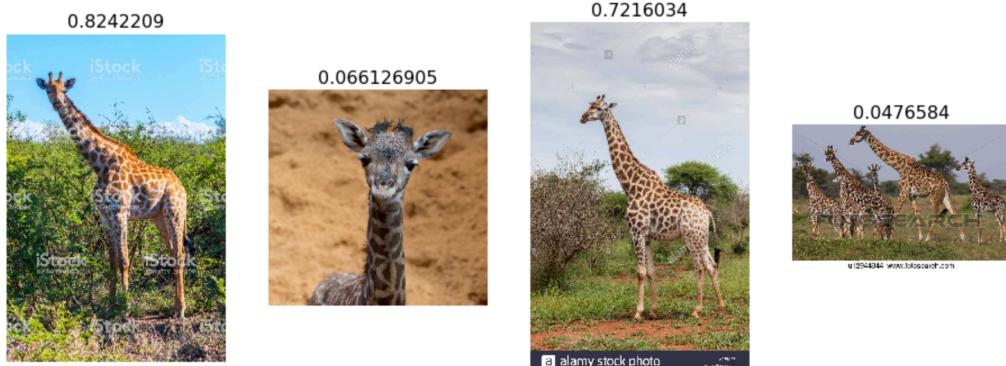
Epoch  100% 10/10 [00:00<00:00, 17.28it/s]

| epoch | trn_loss | val_loss | accuracy |
|-------|----------|----------|----------|
| 0     | 0.755909 | 0.806627 | 0.45     |
| 1     | 0.729105 | 0.759772 | 0.5      |
| 2     | 0.684765 | 0.71894  | 0.575    |
| 3     | 0.633249 | 0.695716 | 0.55     |
| 4     | 0.580498 | 0.718255 | 0.575    |
| 5     | 0.51721  | 0.747415 | 0.55     |
| 6     | 0.478577 | 0.779054 | 0.575    |
| 7     | 0.447871 | 0.822055 | 0.575    |
| 8     | 0.409474 | 0.847211 | 0.575    |
| 9     | 0.376428 | 0.862352 | 0.575    |

```
Out[327]: [array([0.86235]), 0.574999988079071]
```

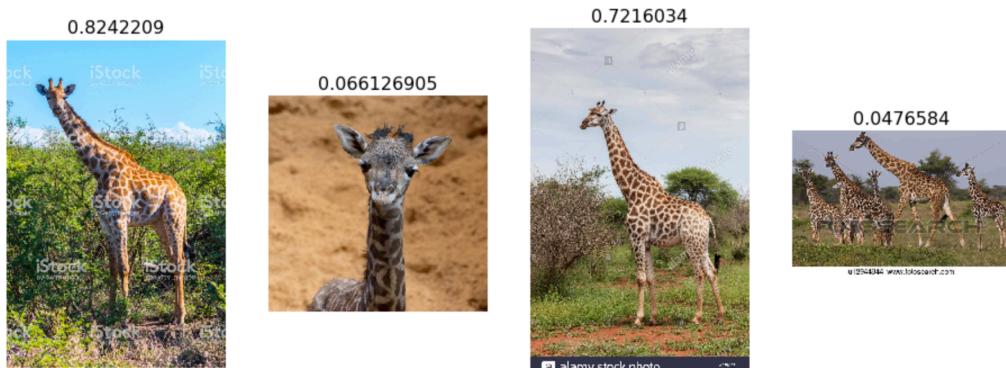
```
In [336]: # 1. A few correct labels at random  
plot_val_with_title(rand_by_correct(True), "Correctly classified")
```

Correctly classified



```
In [336]: # 1. A few correct labels at random  
plot_val_with_title(rand_by_correct(True), "Correctly classified")
```

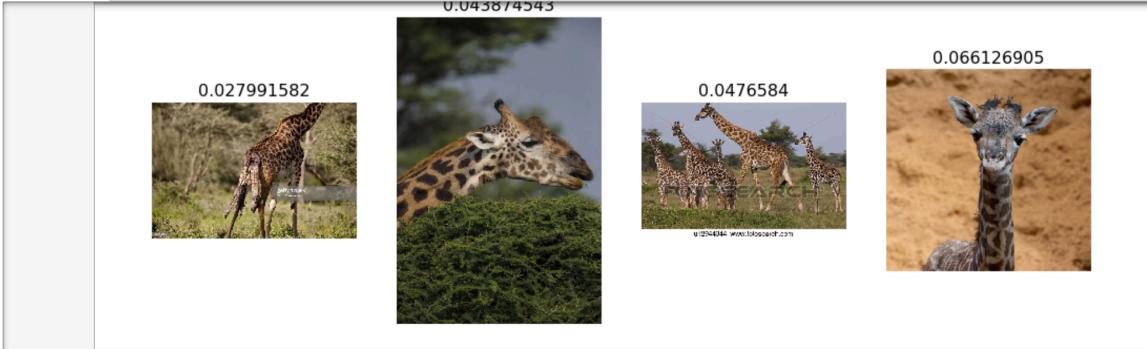
Correctly classified



```
In [337]: # 2. A few incorrect labels at random  
plot_val_with_title(rand_by_correct(False), "Incorrectly classified")  
Incorrectly classified
```



```
In [339]: plot_val_with_title(most_by_correct(0, True), f"Most correct {class1}")
```

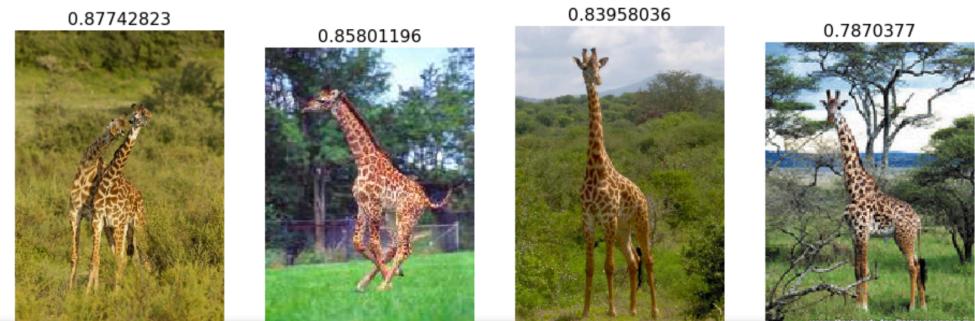


```
In [340]: plot_val_with_title(most_by_correct(1, True), "Most correct {class2}")  
Most correct {class2}
```



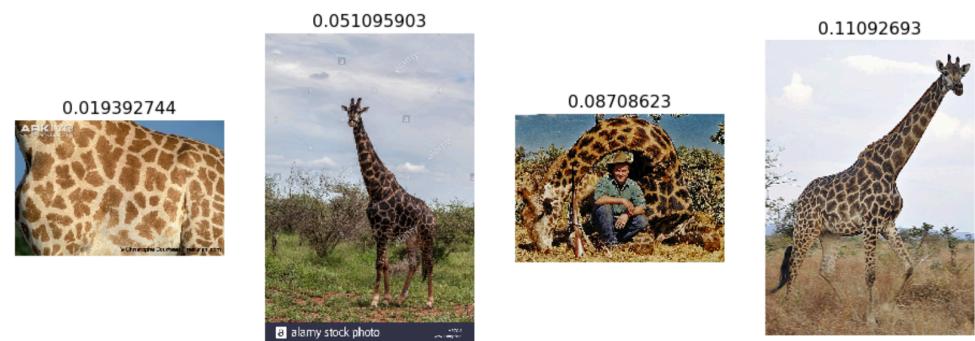
```
In [341]: plot_val_with_title(most_by_correct(0, False), "Most incorrect class 1")
```

Most incorrect class 1



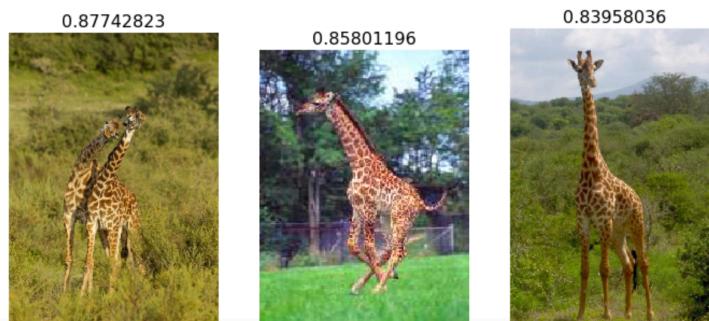
```
# In [342]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")
```

Most incorrect class 2



```
In [341]: plot_val_with_title(most_by_correct(0, False), "Most incorrect class 1")
```

Most incorrect class 1



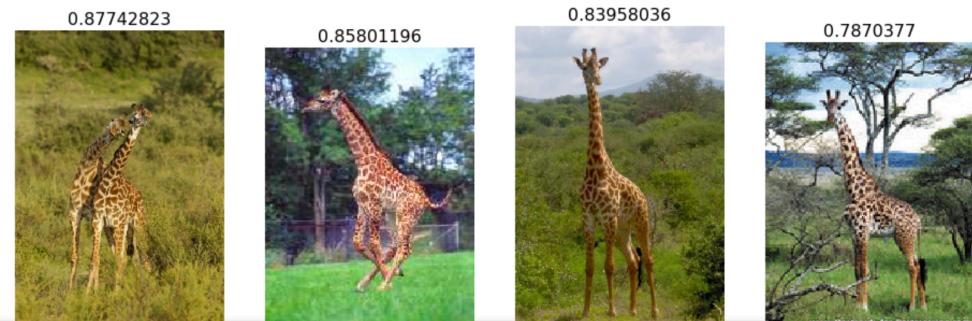
```
# In [342]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")
```

Most incorrect class 2



```
In [341]: plot_val_with_title(most_by_correct(0, False), "Most incorrect class 1")
```

Most incorrect class 1



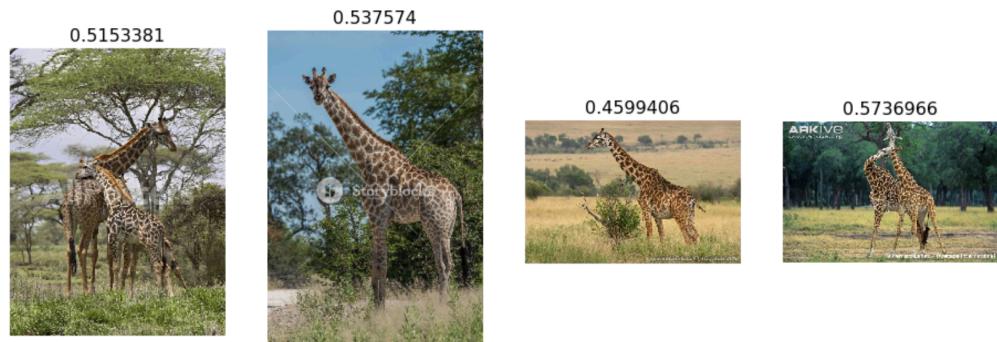
```
In [342]: plot_val_with_title(most_by_correct(1, False), "Most incorrect class 2")
```

Most incorrect class 2



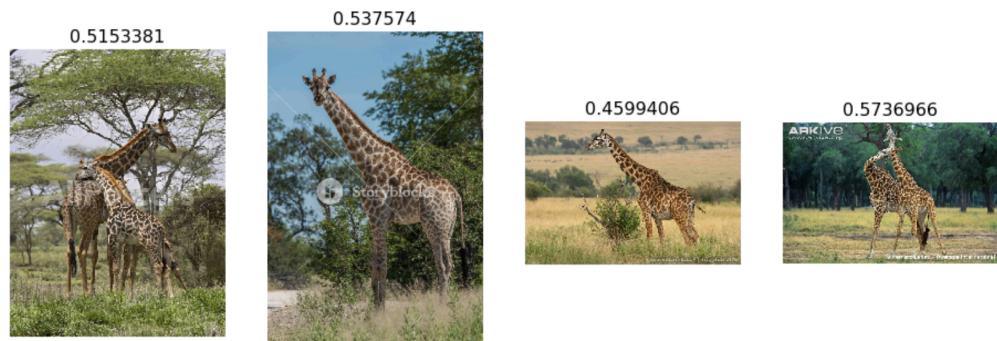
```
In [343]: most_uncertain = np.argsort(np.abs(probs - 0.5))[:4]
plot_val_with_title(most_uncertain, "Most uncertain predictions")
```

Most uncertain predictions



```
In [343]: most_uncertain = np.argsort(np.abs(probs - 0.5))[:4]
plot_val_with_title(most_uncertain, "Most uncertain predictions")
```

Most uncertain predictions



Conclusion: as I expected, a rather hard task. The NN does slightly better than chance. Worse than I predicted. This is an interesting challenge to pick up next. I feel that some tuning of the model, or more data / data augmentation can really help here.

TODO apple versus peach